

```
In [1]: sc
```

```
Out[1]: SparkContext
```

[Spark UI \(http://172.16.86.148:4040\)](http://172.16.86.148:4040)

Version

v2.2.1

Master

local[2]

AppName

pyspark-shell

Read the streams from kafka

```
In [2]: df = spark \
        .readStream \
        .format("kafka") \
        .option("kafka.bootstrap.servers", "localhost:9092") \
        .option("subscribe", "tweets_topic") \
        .load()
```

Possible kafka read options

- Reading from multiple topics

```
option("subscribe", "topicA,topicB")
```

- Reading from topics with names following a pattern

```
option("subscribepattern", ""topic\d"")
```

- Reading topics and partitions

```
assign({"topicA":[0,1],"topicB":[0,1]})
```

- Start reading from which offset

```
option("startingoffsets", "latest") // or "earliest"
```

- How many record for each trigger

```
option("maxOffsetsPerTrigger", 1)
```

Extract tags and count them

- Information that spark reads for each record

```
val result = df.
  select(
    $"key" cast "string", // deserialize keys
    $"value" cast "string", // deserialize values
    $"topic",
    $"partition",
    $"offset")
```

```
In [3]: from pyspark.sql.functions import explode, split, col

words = df.select(explode(split(df.value, " ")).alias("token"))
tags = words.where(words.token.contains("#"))
tagcount = tags.groupBy("token").count()
```

Write results to the sink

```
In [4]: query = tagcount.writeStream \
        .outputMode("complete") \
        .format("memory") \
        .queryName("tweetstrends") \
        .start()
```

- Setting trigger interval

```
trigger(Trigger.ProcessingTime(10.seconds))
```

- Checkpoint path to recover from failures

```
option("checkpointLocation", "path/to/HDFS/dir") \
```

Print the top 10 trends

```
In [9]: ## Show the top 10 tags
spark.sql("Select token, count from tweetstrends order by count desc limit 10").show()
```

```
+-----+-----+
|          token|count|
+-----+-----+
|#INSTANTFOLLOWBACK|    1|
|          #handmade|    1|
|          #sidoarjo|    1|
|          #sweet|    1|
|          #art|    1|
|    #goodmoning|    1|
|    #colazione|    1|
|          #Aries|    1|
|          #final|    1|
|    #TEAMFOLLOWWACK|    1|
+-----+-----+
```

Check Status

```
In [6]: query.lastProgress
```


```
Out[6]: {'id': 'b9d1ad9b-02c0-4f6f-817a-403e7acb2d32',  
        'runId': 'da7147bb-3988-4a9e-b3d2-0a403b2275b7',  
        'name': 'tweetstrends',  
        'timestamp': '2018-08-28T20:14:19.624Z',  
        'numInputRows': 23,  
        'inputRowsPerSecond': 1533.3333333333335,  
        'processedRowsPerSecond': 4.03013842649378,  
        'durationMs': {'addBatch': 5588,  
                        'getBatch': 15,  
                        'getOffset': 1,  
                        'queryPlanning': 31,  
                        'triggerExecution': 5707,  
                        'walCommit': 68},  
        'stateOperators': [{'numRowsTotal': 3, 'numRowsUpdated': 3}],  
        'sources': [{'description': 'KafkaSource[Subscribe[tweets_topic]]',  
                     'startOffset': {'tweets_topic': {'0': 1219}},  
                     'endOffset': {'tweets_topic': {'0': 1242}},  
                     'numInputRows': 23,  
                     'inputRowsPerSecond': 1533.3333333333335,  
                     'processedRowsPerSecond': 4.03013842649378}],  
        'sink': {'description': 'MemorySink'}}
```

Visualize the Trends

```
In [10]: import matplotlib.pyplot as plt  
import seaborn as sn  
%matplotlib inline
```

```
In [11]: import time
        from IPython import display

count = 0
while count < 10:
    time.sleep( 1 )
    top_10_tweets = sqlContext.sql( 'Select token, count from tweetstrends order by count desc limit 10' )
    top_10_df = top_10_tweets.toPandas()
    display.clear_output(wait=True)
    plt.figure( figsize = ( 10, 8 ) )
    sn.barplot( x="count", y="token", data=top_10_df)
    plt.show()
    count = count + 1
```



In []: